



# CODING HORROR

programming and human factors  
by Jeff Atwood

## A Visual Explanation of SQL Joins

October 11, 2007

I thought Ligaya Turmelle's [post on SQL joins](#) was a great primer for novice developers. Since SQL joins *appear* to be set-based, the use of [Venn diagrams](#) to explain them seems, at first blush, to be a natural fit. However, like the commenters to her post, I found that the Venn diagrams didn't quite match the [SQL join syntax](#) reality in my testing.

I love the concept, though, so let's see if we can make it work. Assume we have the following two tables. **Table A** is on the left, and **Table B** is on the right. We'll populate them with four records each.


```
=====
id  name      id  name
--  ----      --  ----
1   Pirate    1   Rutabaga
2   Monkey    2   Pirate
3   Ninja     3   Darth Vader
4   Spaghetti 4   Ninja
=====
```

Let's join these tables by the name field in a few different ways and see if we can get a conceptual match to those nifty Venn diagrams.

```
=====
SELECT * FROM TableA
INNER JOIN TableB
ON TableA.name = TableB.name
=====
```

```
id  name      id  name
--  ----      --  ----
1   Pirate    2   Pirate
3   Ninja     4   Ninja
=====
```

**Inner join** produces only the set of records that match in both Table A and Table B.




```
=====
SELECT * FROM TableA
FULL OUTER JOIN TableB
ON TableA.name = TableB.name
=====
```

```
id  name      id  name
--  ----      --  ----
1   Pirate    2   Pirate
2   Monkey    null null
3   Ninja     4   Ninja
4   Spaghetti null null
null null    1   Rutabaga
null null    3   Darth Vader
=====
```

**Full outer join** produces the set of all records in Table A and Table B, with matching records from both sides


where available. If there is no match, the missing side will contain null.



```
=====
SELECT * FROM TableA
LEFT OUTER JOIN TableB
ON TableA.name = TableB.name
=====
```

| id | name      | id   | name   |
|----|-----------|------|--------|
| -- | ----      | --   | ----   |
| 1  | Pirate    | 2    | Pirate |
| 2  | Monkey    | null | null   |
| 3  | Ninja     | 4    | Ninja  |
| 4  | Spaghetti | null | null   |


**Left outer join** produces a complete set of records from Table A, with the matching records (where available) in Table B. If there is no match, the right side will contain null.



```
=====
SELECT * FROM TableA
LEFT OUTER JOIN TableB
ON TableA.name = TableB.name
WHERE TableB.id IS null
=====
```

| id | name      | id   | name |
|----|-----------|------|------|
| -- | ----      | --   | ---- |
| 2  | Monkey    | null | null |
| 4  | Spaghetti | null | null |

To produce the set of records only in Table A, but not in Table B, we perform the same left outer join, then **exclude the records we don't want from the right side via a where clause.**



```


=====
SELECT * FROM TableA
FULL OUTER JOIN TableB
ON TableA.name = TableB.name
WHERE TableA.id IS null
OR TableB.id IS null
=====

```

| id   | name      | id   | name        |
|------|-----------|------|-------------|
| --   | ----      | --   | ----        |
| 2    | Monkey    | null | null        |
| 4    | Spaghetti | null | null        |
| null | null      | 1    | Rutabaga    |
| null | null      | 3    | Darth Vader |

```
=====
```

To produce the set of records unique to Table A and Table B, we perform the same full outer join, then **exclude the records we don't want from both sides via a where clause.**



There's also a cartesian product or **cross join**, which as far as I can tell, can't be expressed as a Venn diagram:

```

=====
SELECT * FROM TableA
CROSS JOIN TableB
=====

```

This joins "everything to everything", resulting in  $4 \times 4 = 16$  rows, far more than we had in the original sets. If you do the math, you can see why this is a *very* dangerous join to run against large tables.

Posted by Jeff Atwood

#### Related posts:

- [Why Does Windows Have Terrible Battery Life?](#)
- [You Don't Need Millions of Dollars](#)
- [Updating Your Utility Belt](#)
- [The 2013 HTPC Build](#)
- [The CODE Keyboard](#)